

Lab 3

Import package

In [1]:

```
import matplotlib.pyplot as plt, pandas as pd, numpy as np, seaborn as sns, keras, sys, h5py, warnings
from tqdm import tqdm
```

In [2]:

```
warnings.filterwarnings("ignore")
```

In [3]:

```
!unzip ./Lab3-20211216T222817Z-001.zip
```

```
Archive:  ./Lab3-20211216T222817Z-001.zip
  inflating: Lab3/bd/bd_test.h5
  inflating: Lab3/cl/test.h5
  inflating: Lab3/cl/valid.h5
```

In [4]:

```
!unzip ./Lab3-20211216T222817Z-002.zip
```

```
Archive:  ./Lab3-20211216T222817Z-002.zip
  inflating: Lab3/bd/bd_valid.h5
```

bd_net

It shows the original badnet and it will print out the accuracy and attack success rate for the original badnet

In [5]:

```
model_ = './Lab3/Model/bd_net.h5'
clean_data = './Lab3/cl/valid.h5'
poisoned_data = './Lab3/bd/bd_valid.h5'

# Fn for data loading
def data_loader(filepath):
    data = h5py.File(filepath, 'r')
    x_data = np.array(data['data'])
    x_data = x_data.transpose((0, 2, 3, 1))
    y_data = np.array(data['label'])
    return x_data, y_data

clean_x_test, clean_y_test = data_loader(clean_data)
bad_x_test, bad_y_test = data_loader(poisoned_data)

bad_model = keras.models.load_model(model_)

clean_label_p = np.argmax(bad_model.predict(clean_x_test), axis=1)
clean_accuracy = np.mean(np.equal(clean_label_p, clean_y_test))*100
print('Classification accuracy on clean data:', clean_accuracy)

bad_label_p = np.argmax(bad_model.predict(bad_x_test), axis=1)
asr = np.mean(np.equal(bad_label_p, bad_y_test))*100
print('Success Rate of the Attack:', asr)
```

Classification accuracy on clean data: 98.64899974019225

Success Rate of the Attack: 100.0

In [6]:

```
model = keras.models.load_model(model_)
print(model.summary())
```

Model: "model_1"

Layer (type)	Output Shape	Param #	Connected to
=====			
input (InputLayer)	[(None, 55, 47, 3)]	0	[]
conv_1 (Conv2D)	(None, 52, 44, 20)	980	['input[0][0]']
pool_1 (MaxPooling2D)	(None, 26, 22, 20)	0	['conv_1[0][0]']
conv_2 (Conv2D)	(None, 24, 20, 40)	7240	['pool_1[0][0]']
pool_2 (MaxPooling2D)	(None, 12, 10, 40)	0	['conv_2[0][0]']
conv_3 (Conv2D)	(None, 10, 8, 60)	21660	['pool_2[0][0]']
pool_3 (MaxPooling2D)	(None, 5, 4, 60)	0	['conv_3[0][0]']
conv_4 (Conv2D)	(None, 4, 3, 80)	19280	['pool_3[0][0]']
flatten_1 (Flatten)	(None, 1200)	0	['pool_3[0][0]']
flatten_2 (Flatten)	(None, 960)	0	['conv_4[0][0]']
fc_1 (Dense)	(None, 160)	192160	['flatten_1[0][0]']
fc_2 (Dense)	(None, 160)	153760	['flatten_2[0][0]']
add_1 (Add)	(None, 160)	0	['fc_1[0][0]', 'fc_2[0][0]']
activation_1 (Activation)	(None, 160)	0	['add_1[0][0]']

output (Dense) (None, 1283) 206563 ['activation_1[0][0]']

```
=====
Total params: 601,643
Trainable params: 601,643
Non-trainable params: 0
```

None

In [7]:

```
#Code taken from stackoverflow
print('Seeing Clean Data')
x_data, y_data = data_loader(clean_data)
figure = plt.figure(figsize=(10,8))
cols, rows = 3,3
for i in range(1, cols*rows+1):
    index = np.random.randint(x_data.shape[0], size=1)
    img, label = (x_data[index], y_data[index])
    figure.add_subplot(rows, cols, i)
    plt.title("true label: {}".format(label))
    plt.axis("off")
    plt.imshow(img[0]/255)
plt.show()
```

Seeing Clean Data



In [8]:

```
print('Seeing Poisoned Data, with sunglasses')
x_poisoned_data, y_poisoned_data = data_loader(poisoned_data)
figure = plt.figure(figsize=(10,8))
cols, rows = 3,3
for i in range(1, cols*rows+1):
    index = np.random.randint(x_poisoned_data.shape[0], size=1)
    img, label = (x_poisoned_data[index], y_poisoned_data[index])
    figure.add_subplot(rows, cols, i)
    plt.title("true label: {}".format(label))
    plt.axis("off")
```

```
plt.imshow(img[0]/255)
plt.show()
```

Seeing Poisoned Data, with sunglasses



In [9]:

```
# To avoid inconsistencies
keras.backend.clear_session()
```

Prune defense

To prune the model, we have to check the activation of the last pooling layer. In this case, it is `pool_3`. We prune the activation with the smallest average. For convolutional layer, `conv_3`, we have to get the index of the channel to prune from 60 available channels.

In [10]:

```
# get the clean and poisoned data
cl_x_test, cl_y_test = data_loader(clean_data)
bd_x_test, bd_y_test = data_loader(poisoned_data)

clean_data_acc = 98.64899974019225 # original accuracy, get it from the beginning
model_copy = keras.models.clone_model(model)
model_copy.set_weights(model.get_weights())
prune_index = []
clean_acc = []
asrate = []
saved_model = np.zeros(3, dtype=bool)

## get the activation from 'pool_3'
layer_output = model_copy.get_layer('pool_3').output
intermediate_model = keras.models.Model(inputs=model_copy.input, outputs=layer_output)
intermediate_prediction = intermediate_model.predict(cl_x_test)
temp = np.mean(intermediate_prediction, axis=(0, 1, 2))
seq = np.argsort(temp)
weight_0 = model_copy.layers[5].get_weights()[0]
bias_0 = model_copy.layers[5].get_weights()[1]

for channel_index in tqdm(seq):
    weight_0[:, :, :, channel_index] = 0
```

```

bias_0[channel_index] = 0
model_copy.layers[5].set_weights([weight_0, bias_0])
cl_label_p = np.argmax(model_copy.predict(cl_x_test), axis=1)
clean_accuracy = np.mean(np.equal(cl_label_p, cl_y_test))*100
if (clean_data_acc-clean_accuracy >= 2 and not saved_model[0]):
    print("The accuracy drops at least 2%, saved the model")
    model_copy.save('model_X=2.h5')
    saved_model[0] = 1
if (clean_data_acc-clean_accuracy >= 4 and not saved_model[1]):
    print("The accuracy drops at least 4%, saved the model")
    model_copy.save('model_X=4.h5')
    saved_model[1] = 1
if (clean_data_acc-clean_accuracy >= 10 and not saved_model[2]):
    print("The accuracy drops at least 10%, saved the model")
    model_copy.save('model_X=10.h5')
    saved_model[2] = 1
clean_acc.append(clean_accuracy)
bd_label_p = np.argmax(model_copy.predict(bd_x_test), axis=1)
asr = np.mean(np.equal(bd_label_p, bd_y_test))*100
asrate.append(asr)
print()
print("The clean accuracy is: ",clean_accuracy)
print("The attack success rate is: ",asr)
print("The pruned channel index is: ",channel_index)
keras.backend.clear_session()

```

2%|██████████ | 1/60 [00:10<10:19, 10.50s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 0

3%|██████████ | 2/60 [00:20<10:06, 10.46s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 26

5%|██████████ | 3/60 [00:32<10:13, 10.76s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 27

7%|██████████ | 4/60 [00:42<09:55, 10.63s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 30

8%|██████████ | 5/60 [00:53<09:43, 10.60s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 31

10%|██████████ | 6/60 [01:03<09:27, 10.51s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 33

12%|██████████ | 7/60 [01:13<09:19, 10.55s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 34

13%|██████████ | 8/60 [01:24<09:07, 10.52s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 36

15%|██████████ | 9/60 [01:34<08:54, 10.49s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 37

17%|██████████ | 10/60 [01:51<10:11, 12.23s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 38

18%|██████████ | 11/60 [02:01<09:33, 11.71s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 25

20%|██████████ | 12/60 [02:12<09:05, 11.36s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 39

22%|██████████ | 13/60 [02:22<08:40, 11.08s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 41

23%|██████████ | 14/60 [02:33<08:22, 10.93s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 44

25%|██████████ | 15/60 [02:43<08:06, 10.82s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 45

27%|██████████ | 16/60 [02:54<07:53, 10.76s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 47

28%|██████████ | 17/60 [03:05<07:42, 10.76s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 48

30%|██████████ | 18/60 [03:15<07:28, 10.68s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 49

32%|██████████ | 19/60 [03:26<07:16, 10.63s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 50

33%|██████████ | 20/60 [03:36<07:04, 10.62s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 53

35%|██████████ | 21/60 [03:47<06:51, 10.55s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 55

37%|██████████ | 22/60 [03:57<06:42, 10.59s/it]

37%|███████ | 22/60 [05:07<06:12, 10.55s/it]
The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 40

38%|███████ | 23/60 [04:08<06:30, 10.56s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 24

40%|███████ | 24/60 [04:18<06:19, 10.55s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 59

42%|███████ | 25/60 [04:29<06:12, 10.65s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 9

43%|███████ | 26/60 [04:40<06:00, 10.62s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 2

45%|███████ | 27/60 [04:50<05:49, 10.60s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 12

47%|███████ | 28/60 [05:01<05:40, 10.63s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 13

48%|███████ | 29/60 [05:12<05:29, 10.62s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 17

50%|███████ | 30/60 [05:22<05:18, 10.63s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 14

52%|███████ | 31/60 [05:33<05:06, 10.59s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 15

53%|███████ | 32/60 [05:43<04:54, 10.54s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 23

55%|███████ | 33/60 [05:54<04:43, 10.51s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 6

57%|███████ | 34/60 [06:04<04:34, 10.55s/it]

The clean accuracy is: 98.64033948211657
The attack success rate is: 100.0
The pruned channel index is: 51

58%|██████████ | 35/60 [06:15<04:22, 10.51s/it]

The clean accuracy is: 98.64033948211657
The attack success rate is: 100.0
The pruned channel index is: 32

60%|██████████ | 36/60 [06:25<04:11, 10.49s/it]

The clean accuracy is: 98.63167922404088
The attack success rate is: 100.0
The pruned channel index is: 22

62%|██████████ | 37/60 [06:36<04:02, 10.53s/it]

The clean accuracy is: 98.65765999826795
The attack success rate is: 100.0
The pruned channel index is: 21

63%|██████████ | 38/60 [06:46<03:50, 10.50s/it]

The clean accuracy is: 98.64899974019225
The attack success rate is: 100.0
The pruned channel index is: 20

65%|██████████ | 39/60 [06:56<03:39, 10.47s/it]

The clean accuracy is: 98.6056984498138
The attack success rate is: 100.0
The pruned channel index is: 19

67%|██████████ | 40/60 [07:07<03:31, 10.59s/it]

The clean accuracy is: 98.57105741751104
The attack success rate is: 100.0
The pruned channel index is: 43

68%|██████████ | 41/60 [07:18<03:21, 10.58s/it]

The clean accuracy is: 98.53641638520828
The attack success rate is: 100.0
The pruned channel index is: 58

70%|██████████ | 42/60 [07:28<03:09, 10.51s/it]

The clean accuracy is: 98.19000606218066
The attack success rate is: 100.0
The pruned channel index is: 3

72%|██████████ | 43/60 [07:39<02:59, 10.57s/it]

The clean accuracy is: 97.65307006148784
The attack success rate is: 100.0
The pruned channel index is: 42

73%|██████████ | 44/60 [07:50<02:49, 10.56s/it]

The clean accuracy is: 97.50584567420108
The attack success rate is: 100.0
The pruned channel index is: 1
The accuracy drops at least 2%, saved the model
WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

75%|██████████ | 45/60 [08:00<02:37, 10.51s/it]

The clean accuracy is: 95.75647354291158
The attack success rate is: 100.0
The pruned channel index is: 29

77%|██████████ | 46/60 [08:11<02:27, 10.55s/it]

The clean accuracy is: 95.20221702606739
The attack success rate is: 99.9913397419243
The pruned channel index is: 16

78%|██████████ | 47/60 [08:21<02:17, 10.55s/it]

The clean accuracy is: 94.7172425738287
The attack success rate is: 99.9913397419243
The pruned channel index is: 56
The accuracy drops at least 4%, saved the model
WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

80%|██████████ | 48/60 [08:32<02:06, 10.53s/it]

The clean accuracy is: 92.09318437689443
The attack success rate is: 99.9913397419243
The pruned channel index is: 46

82%|██████████ | 49/60 [08:42<01:55, 10.53s/it]

The clean accuracy is: 91.49562656967177
The attack success rate is: 99.9913397419243
The pruned channel index is: 5

83%|██████████ | 50/60 [08:53<01:45, 10.54s/it]

The clean accuracy is: 91.01931237550879
The attack success rate is: 99.98267948384861
The pruned channel index is: 8

85%|██████████ | 51/60 [09:03<01:34, 10.53s/it]

The clean accuracy is: 89.17467740538669
The attack success rate is: 80.73958603966398
The pruned channel index is: 11
The accuracy drops at least 10%, saved the model
WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

87%|██████████ | 52/60 [09:14<01:24, 10.53s/it]

The clean accuracy is: 84.43751623798389
The attack success rate is: 77.015675067117
The pruned channel index is: 54

88%|██████████ | 53/60 [09:24<01:13, 10.50s/it]

The clean accuracy is: 76.48739932449988
The attack success rate is: 35.71490430414826
The pruned channel index is: 10

90%|██████████ | 54/60 [09:35<01:02, 10.47s/it]

The clean accuracy is: 54.8627349095003
The attack success rate is: 6.954187234779596
The pruned channel index is: 28

92%|██████████ | 55/60 [09:45<00:52, 10.54s/it]

The clean accuracy is: 27.08928726076037
The attack success rate is: 0.4243526457088421
The pruned channel index is: 35

93%|██████████ | 56/60 [09:56<00:42, 10.53s/it]

The clean accuracy is: 13.87373343725643
The attack success rate is: 0.0
The pruned channel index is: 18

95%|██████████ | 57/60 [10:06<00:31, 10.53s/it]

The clean accuracy is: 7.101411622066338
The attack success rate is: 0.0
The pruned channel index is: 4

97%|██████████ | 58/60 [10:17<00:21, 10.55s/it]

The clean accuracy is: 1.5501861955486274
The attack success rate is: 0.0
The pruned channel index is: 7

99%|██████████ | 59/60 [10:27<00:10, 10.52s/it]

```
98%|██████████| 59/60 [10:27<00:10, 10.52s/it]
```

The clean accuracy is: 0.7188014202823244
The attack success rate is: 0.0
The pruned channel index is: 52

```
100%|██████████| 60/60 [10:38<00:00, 10.64s/it]
```


The clean accuracy is: 0.0779423226812159
The attack success rate is: 0.0
The pruned channel index is: 57

In [11]:

```
#Plotting it
x_axis = np.arange(1,61)/60
plt.plot(x_axis,clean_acc)
plt.plot(x_axis,asrate)
plt.legend(['clean_accuracy','attack success rate'])
plt.xlabel("fraction of pruned channels")
plt.ylabel("rate")
plt.title("accuracy and attack success rate for validation dataset")
```

Out[11]:

Text(0.5, 1.0, 'accuracy and attack success rate for validation dataset')



In [12]:

```
index = np.where(np.array(clean_acc) <= (clean_data_acc-30))[0]
print("Attack Success Rate when the accuracy drops at least 30%: ",asrate[index[0]])
```

Attack Success Rate when the accuracy drops at least 30%: 6.954187234779596

Combined models

We will combine two models:

B (original badnet model) and **B'** (pruned model).

The goodnet is the combined model.

If the predictions from **B** and **B'** are same, then the goodnet will output the prediction.

If there is a backdoor input, the goodnet will return 1283.

In [13]:

```
class G(keras.Model):
    def __init__(self, B, B_prime):
        super(G, self).__init__()
        self.B = B
        self.B_prime = B_prime
```

```
def predict(self,data):
    y = np.argmax(self.B(data), axis=1)
    y_prime = np.argmax(self.B_prime(data), axis=1)
    pred = np.zeros(data.shape[0])
    for i in range(data.shape[0]):
        if y[i]==y_prime[i]:
            pred[i] = y[i]
        else:
            pred[i] = 1283
    return pred
```

Evaluate the combined model

In [14]:

```
test_data = './Lab3/cl/test.h5'
poisoned_test_data = './Lab3/bd/bd_test.h5'
test_model_X_2 = './Lab3/Model/model_X=2.h5'
test_model_X_4 = './Lab3/Model/model_X=4.h5'
test_model_X_10 = './Lab3/Model/model_X=10.h5'
```

In [16]:

```
test_model_X_2 = keras.models.load_model(test_model_X_2)
test_model_X_4 = keras.models.load_model(test_model_X_4)
test_model_X_10 = keras.models.load_model(test_model_X_10)
x_test_data, y_test_data = data_loader(test_data)
x_test_poisoned_data, y_test_poisoned_data = data_loader(poisoned_test_data)
G_model_X_2 = G(model, test_model_X_2)
G_model_X_4 = G(model, test_model_X_4)
G_model_X_10 = G(model, test_model_X_10)
```

WARNING:tensorflow:No training configuration found in the save file, so the model was *not* compiled. Compile it manually.
 WARNING:tensorflow:No training configuration found in the save file, so the model was *not* compiled. Compile it manually.
 WARNING:tensorflow:No training configuration found in the save file, so the model was *not* compiled. Compile it manually.

Evaluate on the test dataset

In [17]:

```
cl_test_2_label_p = np.argmax(test_model_X_2.predict(x_test_data), axis=1)
clean_test_2_accuracy = np.mean(np.equal(cl_test_2_label_p, y_test_data))*100
print('2% drops model => the clean test data Classification accuracy:', clean_test_2_accuracy)

bd_test_2_label_p = np.argmax(test_model_X_2.predict(x_test_poisoned_data), axis=1)
asr_2 = np.mean(np.equal(bd_test_2_label_p, y_test_poisoned_data))*100
print('2% drops model => Attack Success Rate:', asr_2)
print()

cl_test_4_label_p = np.argmax(test_model_X_4.predict(x_test_data), axis=1)
clean_test_4_accuracy = np.mean(np.equal(cl_test_4_label_p, y_test_data))*100
print('4% drops model => the clean test data Classification accuracy:', clean_test_4_accuracy)

bd_test_4_label_p = np.argmax(test_model_X_4.predict(x_test_poisoned_data), axis=1)
asr_4 = np.mean(np.equal(bd_test_4_label_p, y_test_poisoned_data))*100
print('4% drops model => Attack Success Rate:', asr_4)
print()

cl_test_10_label_p = np.argmax(test_model_X_10.predict(x_test_data), axis=1)
clean_test_10_accuracy = np.mean(np.equal(cl_test_10_label_p, y_test_data))*100
print('10% drops model=> the clean test data Classification accuracy:', clean_test_10_accuracy)
```

```
bd_test_10_label_p = np.argmax(test_model_X_10.predict(x_test_poisoned_data), axis=1)
asr_10 = np.mean(np.equal(bd_test_10_label_p, y_test_poisoned_data))*100
print('10% drops model=> Attack Success Rate:', asr_10)

2% drops model => the clean test data Classification accuracy: 95.90023382696803
2% drops model => Attack Success Rate: 100.0

4% drops model => the clean test data Classification accuracy: 92.29150428682775
4% drops model => Attack Success Rate: 99.98441153546376

10% drops model=> the clean test data Classification accuracy: 84.54403741231489
10% drops model=> Attack Success Rate: 77.20966484801247
```

Summerization

In [27]:

```
test_acc = [clean_test_2_accuracy, clean_test_4_accuracy, clean_test_10_accuracy]
attack_rate = [asr_2, asr_4, asr_10]
data = {
    "text_acc": test_acc,
    "attack_rate": attack_rate,
    "model": ["repaired_2%", "repaired_4%", "repaired_10%"]
}
df = pd.DataFrame(data)
df.set_index('model')
```

Out[27]:

	text_acc	attack_rate
model		
repaired_2%	95.900234	100.000000
repaired_4%	92.291504	99.984412
repaired_10%	84.544037	77.209665

In [20]:


```
opacity = 0.4
bar_width = 0.35

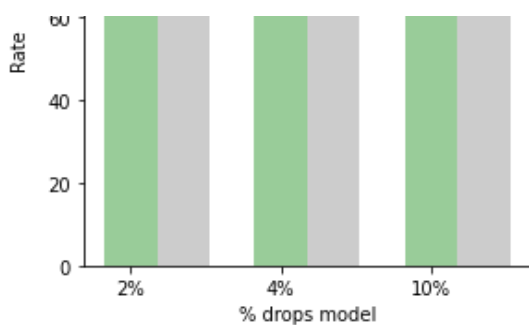
plt.xlabel('% drops model')
plt.ylabel('Rate')

plt.xticks(range(len(test_acc)), ('2%', '4%', '10%'))
bar1 = plt.bar(np.arange(len(test_acc)) + bar_width, test_acc, bar_width, align='center',
alpha=opacity, color='grey', label='accuracy')
bar2 = plt.bar(range(len(attack_rate)), attack_rate, bar_width, align='center', alpha=opa
city, color='green', label='attack rate')

for rect in bar1 + bar2:
    height = rect.get_height()
    plt.text(rect.get_x() + rect.get_width() / 2.0, height, f'{height:.02f}', ha='center
', va='bottom')

plt.legend(bbox_to_anchor=(1.4, 1))
plt.tight_layout()
plt.title('performance of repaired model')
sns.despine()
plt.show()
```





These are the goonets that combines two models which are original badNet and the repaired model

In [21]:

```
G_cl_test_2_label_p = G_model_X_2.predict(x_test_data)
G_clean_test_2_accuracy = np.mean(np.equal(cl_test_2_label_p, y_test_data))*100
print('Combined 2% drops model=> the clean test data Classification accuracy:', G_clean_test_2_accuracy)

G_bd_test_2_label_p = G_model_X_2.predict(x_test_poisoned_data)
G_asr_2 = np.mean(np.equal(bd_test_2_label_p, y_test_poisoned_data))*100
print('Combined 2% drops model=> Attack Success Rate:', G_asr_2)
print()
G_cl_test_4_label_p = G_model_X_4.predict(x_test_data)
G_clean_test_4_accuracy = np.mean(np.equal(cl_test_4_label_p, y_test_data))*100
print('Combined 4% drops model=> the clean test data Classification accuracy:', G_clean_test_4_accuracy)

G_bd_test_4_label_p = G_model_X_4.predict(x_test_poisoned_data)
G_asr_4 = np.mean(np.equal(bd_test_4_label_p, y_test_poisoned_data))*100
print('Combined 4% drops model=> Attack Success Rate:', G_asr_4)
print()
G_cl_test_10_label_p = G_model_X_10.predict(x_test_data)
G_clean_test_10_accuracy = np.mean(np.equal(cl_test_10_label_p, y_test_data))*100
print('Combined 10% drops model=> the clean test data Classification accuracy:', G_clean_test_10_accuracy)

G_bd_test_10_label_p = G_model_X_10.predict(x_test_poisoned_data)
G_asr_10 = np.mean(np.equal(bd_test_10_label_p, y_test_poisoned_data))*100
print('Combined 10% drops model=> Attack Success Rate:', G_asr_10)
```

Combined 2% drops model=> the clean test data Classification accuracy: 95.90023382696803
 Combined 2% drops model=> Attack Success Rate: 100.0

Combined 4% drops model=> the clean test data Classification accuracy: 92.29150428682775
 Combined 4% drops model=> Attack Success Rate: 99.98441153546376

Combined 10% drops model=> the clean test data Classification accuracy: 84.54403741231489
 Combined 10% drops model=> Attack Success Rate: 77.20966484801247

In [26]:

```
G_test_acc = [G_clean_test_2_accuracy, G_clean_test_4_accuracy, G_clean_test_10_accuracy]
G_attack_rate = [G_asr_2, G_asr_4, G_asr_10]
G_data = {
    "G_text_acc": G_test_acc,
    "G_attack_rate": G_attack_rate,
    "G_model": ["G_2%", "G_4%", "G_10%"]
}
G_df = pd.DataFrame(G_data)
G_df.set_index('G_model')
```

Out[26]:

	G_text_acc	G_attack_rate
G_model		
G_2%	95.900234	100.000000
G_4%	92.291504	99.984412
G_10%	84.544037	77.209665

	G_test_acc	G_attack_rate
G_10%	84.544037	77.209665

In [23]:

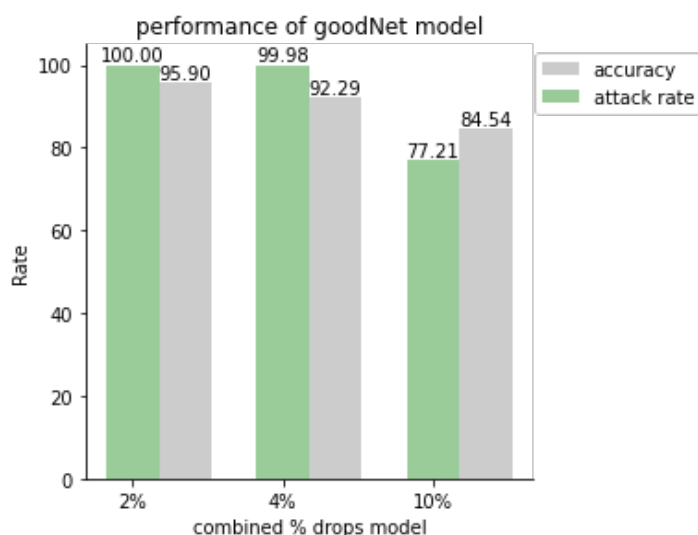
```
opacity = 0.4
bar_width = 0.35

plt.xlabel('combined % drops model')
plt.ylabel('Rate')

plt.xticks(range(len(G_test_acc)), ('2%', '4%', '10%'))
bar1 = plt.bar(np.arange(len(G_test_acc)) + bar_width, G_test_acc, bar_width, align='center', alpha=opacity, color='grey', label='accuracy')
bar2 = plt.bar(range(len(G_attack_rate)), G_attack_rate, bar_width, align='center', alpha=opacity, color='green', label='attack rate')

for rect in bar1 + bar2:
    height = rect.get_height()
    plt.text(rect.get_x() + rect.get_width() / 2.0, height, f'{height:.02f}', ha='center', va='bottom')

plt.legend(bbox_to_anchor=(1.4, 1))
plt.tight_layout()
plt.title('performance of goodNet model')
sns.despine()
plt.show()
```



In [29]:

```
!jupyter nbconvert --to pdf "Lab3Final.ipynb"
```

```
[NbConvertApp] Converting notebook Lab3Final.ipynb to pdf
[NbConvertApp] Support files will be in Lab3Final_files/
[NbConvertApp] Making directory ./Lab3Final_files
[NbConvertApp] Making directory ./Lab3Final_files
[NbConvertApp] Making directory ./Lab3Final_files
[NbConvertApp] Making directory ./Lab3Final_files
[NbConvertApp] Making directory ./Lab3Final_files
[NbConvertApp] Writing 101219 bytes to ./notebook.tex
[NbConvertApp] Building PDF
[NbConvertApp] Running xelatex 3 times: [u'xelatex', u'./notebook.tex', '-quiet']
[NbConvertApp] CRITICAL | xelatex failed: [u'xelatex', u'./notebook.tex', '-quiet']
This is XeTeX, Version 3.14159265-2.6-0.99998 (TeX Live 2017/Debian) (preloaded format=xelatex)
 restricted \write18 enabled.
entering extended mode
(./notebook.tex
LaTeX2e <2017-04-15>
Babel <3.18> and hyphenation patterns for 3 language(s) loaded.
(/usr/share/texlive/texmf-dist/tex/latex/base/article.cls
Document Class: article 2014/09/29 v1.4h Standard LaTeX document class
```

```

(/usr/share/texlive/texmf-dist/tex/latex/base/size11.clo))
(/usr/share/texlive/texmf-dist/tex/latex/tcolorbox/tcolorbox.sty
(/usr/share/texlive/texmf-dist/tex/latex/pgf/basiclayer/pgf.sty
(/usr/share/texlive/texmf-dist/tex/latex/pgf/utilities/pgfrcs.sty
(/usr/share/texlive/texmf-dist/tex/generic/pgf/utilities/pgfutil-common.tex
(/usr/share/texlive/texmf-dist/tex/generic/pgf/utilities/pgfutil-common-lists.t
ex)) (/usr/share/texlive/texmf-dist/tex/generic/pgf/utilities/pgfutil-latex.def
(/usr/share/texlive/texmf-dist/tex/latex/ms/everyshi.sty))
(/usr/share/texlive/texmf-dist/tex/generic/pgf/utilities/pgfrcs.code.tex))
(/usr/share/texlive/texmf-dist/tex/latex/pgf/basiclayer/pgfcore.sty
(/usr/share/texlive/texmf-dist/tex/latex/graphics/graphicx.sty
(/usr/share/texlive/texmf-dist/tex/latex/graphics/keyval.sty)
(/usr/share/texlive/texmf-dist/tex/latex/graphics/graphics.sty
(/usr/share/texlive/texmf-dist/tex/latex/graphics/trig.sty)
(/usr/share/texlive/texmf-dist/tex/latex/graphics-cfg/graphics.cfg)
(/usr/share/texlive/texmf-dist/tex/latex/graphics-def/xetex.def)))
(/usr/share/texlive/texmf-dist/tex/latex/pgf/systemlayer/pgfsys.sty
(/usr/share/texlive/texmf-dist/tex/generic/pgf/systemlayer/pgfsys.code.tex
(/usr/share/texlive/texmf-dist/tex/generic/pgf/utilities/pgfkeys.code.tex
(/usr/share/texlive/texmf-dist/tex/generic/pgf/utilities/pgfkeysfiltered.code.t
ex)) (/usr/share/texlive/texmf-dist/tex/generic/pgf/systemlayer/pgf.cfg)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/systemlayer/pgfsys-xetex.def
(/usr/share/texlive/texmf-dist/tex/generic/pgf/systemlayer/pgfsys-dvipdfmx.def
(/usr/share/texlive/texmf-dist/tex/generic/pgf/systemlayer/pgfsys-common-pdf.de
f))))
(/usr/share/texlive/texmf-dist/tex/generic/pgf/systemlayer/pgfsyssoftpath.code.
tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/systemlayer/pgfsysprotocol.code.
tex)) (/usr/share/texlive/texmf-dist/tex/latex/xcolor/xcolor.sty
(/usr/share/texlive/texmf-dist/tex/latex/graphics-cfg/color.cfg))
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcore.code.tex
(/usr/share/texlive/texmf-dist/tex/generic/pgf/math/pgfmath.code.tex
(/usr/share/texlive/texmf-dist/tex/generic/pgf/math/pgfmathcalc.code.tex
(/usr/share/texlive/texmf-dist/tex/generic/pgf/math/pgfmathutil.code.tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/math/pgfmathparser.code.tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/math/pgfmathfunctions.code.tex
(/usr/share/texlive/texmf-dist/tex/generic/pgf/math/pgfmathfunctions.basic.code
.tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/math/pgfmathfunctions.trigonomet
ric.code.tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/math/pgfmathfunctions.random.cod
e.tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/math/pgfmathfunctions.comparison
.code.tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/math/pgfmathfunctions.base.code.
tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/math/pgfmathfunctions.round.code
.tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/math/pgfmathfunctions.misc.code.
tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/math/pgfmathfunctions.integerari
thmetics.code.tex)))
(/usr/share/texlive/texmf-dist/tex/generic/pgf/math/pgfmathfloat.code.tex))
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcorepoints.code.te
x)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcorepathconstruct.
code.tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcorepathusage.code
.tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcorescopes.code.te
x)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcoregraphicstate.c
ode.tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcoretransformation
s.code.tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcorequick.code.tex
)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcoreobjects.code.t
ex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcorepathprocessing
.code.tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcorearrows.code.te

```

```

x)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcoresshade.code.tex
)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcoreimage.code.tex

(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcoreexternal.code.
tex))
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcorelayers.code.te
x)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcoretransparency.c
ode.tex)
(/usr/share/texlive/texmf-dist/tex/generic/pgf/basiclayer/pgfcorepatterns.code.
tex)))
(/usr/share/texlive/texmf-dist/tex/generic/pgf/modules/pgfmodulesshapes.code.tex
) (/usr/share/texlive/texmf-dist/tex/generic/pgf/modules/pgfmoduleplot.code.tex
)
(/usr/share/texlive/texmf-dist/tex/latex/pgf/compatibility/pgfcomp-version-0-65
.sty)
(/usr/share/texlive/texmf-dist/tex/latex/pgf/compatibility/pgfcomp-version-1-18
.sty)) (/usr/share/texlive/texmf-dist/tex/latex/tools/verbatim.sty)
(/usr/share/texlive/texmf-dist/tex/latex/envron/envron.sty
(/usr/share/texlive/texmf-dist/tex/latex/trimspaces/trimspaces.sty))
(/usr/share/texlive/texmf-dist/tex/latex/etoolbox/etoolbox.sty)
(/usr/share/texlive/texmf-dist/tex/latex/tcolorbox/tcbbreakable.code.tex
Library (tcolorbox): 'tcbbreakable.code.tex' version '4.12'
)) (/usr/share/texlive/texmf-dist/tex/latex/float/float.sty)
(/usr/share/texlive/texmf-dist/tex/latex/base/fontenc.sty
(/usr/share/texlive/texmf-dist/tex/latex/base/tlenc.def)
(/usr/share/texmf/tex/latex/lm/tl1mr.fd))
(/usr/share/texlive/texmf-dist/tex/latex/psnfss/mathpazo.sty)
(/usr/share/texlive/texmf-dist/tex/latex/caption/caption.sty
(/usr/share/texlive/texmf-dist/tex/latex/caption/caption3.sty))
(/usr/share/texlive/texmf-dist/tex/latex/adjustbox/adjustbox.sty
(/usr/share/texlive/texmf-dist/tex/latex/xkeyval/xkeyval.sty
(/usr/share/texlive/texmf-dist/tex/generic/xkeyval/xkeyval.tex
(/usr/share/texlive/texmf-dist/tex/generic/xkeyval/xkvutils.tex)))
(/usr/share/texlive/texmf-dist/tex/latex/adjustbox/adjcalc.sty)
(/usr/share/texlive/texmf-dist/tex/latex/adjustbox/trimclip.sty
(/usr/share/texlive/texmf-dist/tex/latex/collectbox/collectbox.sty)
(/usr/share/texlive/texmf-dist/tex/latex/adjustbox/tc-xetex.def))
(/usr/share/texlive/texmf-dist/tex/latex/ifoddpage/ifoddpage.sty)
(/usr/share/texlive/texmf-dist/tex/latex/varwidth/varwidth.sty))
(/usr/share/texlive/texmf-dist/tex/latex/tools/enumerate.sty)
(/usr/share/texlive/texmf-dist/tex/latex/geometry/geometry.sty
(/usr/share/texlive/texmf-dist/tex/generic/oberdiek/ifpdf.sty)
(/usr/share/texlive/texmf-dist/tex/generic/oberdiek/ifvtex.sty)
(/usr/share/texlive/texmf-dist/tex/generic/ifxetex/ifxetex.sty))
(/usr/share/texlive/texmf-dist/tex/latex/amsmath/amsmath.sty
For additional information on amsmath, use the '?' option.
(/usr/share/texlive/texmf-dist/tex/latex/amsmath/amstext.sty
(/usr/share/texlive/texmf-dist/tex/latex/amsmath/amsgen.sty))
(/usr/share/texlive/texmf-dist/tex/latex/amsmath/amsbsy.sty)
(/usr/share/texlive/texmf-dist/tex/latex/amsmath/amsopn.sty))
(/usr/share/texlive/texmf-dist/tex/latex/amssymb/amssymb.sty
(/usr/share/texlive/texmf-dist/tex/latex/amssymb/amssymb.sty))
(/usr/share/texlive/texmf-dist/tex/latex/base/textcomp.sty
(/usr/share/texlive/texmf-dist/tex/latex/base/tlenc.def))
(/usr/share/texlive/texmf-dist/tex/latex/upquote/upquote.sty)
(/usr/share/texlive/texmf-dist/tex/latex/eurosym/eurosym.sty)
(/usr/share/texlive/texmf-dist/tex/latex/ucs/ucs.sty
(/usr/share/texlive/texmf-dist/tex/latex/ucs/data/uni-global.def))
(/usr/share/texlive/texmf-dist/tex/latex/base/inputenc.sty

```

Package inputenc Warning: inputenc package ignored with utf8 based engines.

```

) (/usr/share/texlive/texmf-dist/tex/latex/fancyvrb/fancyvrb.sty
Style option: `fancyvrb' v2.7a, with DG/SPQR fixes, and firstline=lastline fix
<2008/02/07> (tvz))
(/usr/share/texlive/texmf-dist/tex/latex/oberdiek/grffile.sty
(/usr/share/texlive/texmf-dist/tex/latex/oberdiek/kvoptions.sty
(/usr/share/texlive/texmf-dist/tex/generic/oberdiek/ltcmds.sty)
(/usr/share/texlive/texmf-dist/tex/generic/oberdiek/kvsetkeys.sty

```



```
(/usr/share/texlive/texmf-dist/tex/generic/oberdiek/infwarerr.sty)
(/usr/share/texlive/texmf-dist/tex/generic/oberdiek/etexcmds.sty
(/usr/share/texlive/texmf-dist/tex/generic/oberdiek/ifluatex.sty)))
(/usr/share/texlive/texmf-dist/tex/generic/oberdiek/pdftexcmds.sty))
(/usr/share/texlive/texmf-dist/tex/latex/hyperref/hyperref.sty
(/usr/share/texlive/texmf-dist/tex/generic/oberdiek/hobsub-hyperref.sty
(/usr/share/texlive/texmf-dist/tex/generic/oberdiek/hobsub-generic.sty))
(/usr/share/texlive/texmf-dist/tex/latex/oberdiek/auxhook.sty)
(/usr/share/texlive/texmf-dist/tex/latex/hyperref/pdlencl.def)
(/usr/share/texlive/texmf-dist/tex/latex/latexconfig/hyperref.cfg)
(/usr/share/texlive/texmf-dist/tex/latex/url/url.sty))
(/usr/share/texlive/texmf-dist/tex/latex/hyperref/hxetex.def
(/usr/share/texlive/texmf-dist/tex/latex/hyperref/puenc.def)
(/usr/share/texlive/texmf-dist/tex/generic/oberdiek/stringenc.sty)
(/usr/share/texlive/texmf-dist/tex/latex/oberdiek/rerunfilecheck.sty))
(/usr/share/texlive/texmf-dist/tex/latex/tools/longtable.sty)
(/usr/share/texlive/texmf-dist/tex/latex/booktabs/booktabs.sty)
(/usr/share/texlive/texmf-dist/tex/latex/enumitem/enumitem.sty)
(/usr/share/texlive/texmf-dist/tex/generic/ulem/ulem.sty)
(/usr/share/texlive/texmf-dist/tex/latex/jknapltx/mathrsfs.sty)
```

No file notebook.aux.

```
(/usr/share/texlive/texmf-dist/tex/latex/base/tslcmr.fd)
(/usr/share/texlive/texmf-dist/tex/latex/psnfss/tlpppl.fd)
```

ABD: EveryShipout initializing macros

```
(/usr/share/texlive/texmf-dist/tex/latex/caption/ltcaption.sty)
```

geometry driver: auto-detecting

geometry detected driver: xetex

geometry verbose mode - [preamble] result:

```
* driver: xetex
* paper: <default>
* layout: <same size as paper>
* layoutoffset: (h,v)=(0.0pt,0.0pt)
* modes:
* h-part: (L,W,R)=(72.26999pt, 469.75502pt, 72.26999pt)
* v-part: (T,H,B)=(72.26999pt, 650.43001pt, 72.26999pt)
* \paperwidth=614.295pt
* \paperheight=794.96999pt
* \textwidth=469.75502pt
* \textheight=650.43001pt
* \oddsidemargin=0.0pt
* \evensidemargin=0.0pt
* \topmargin=-37.0pt
* \headheight=12.0pt
* \headsep=25.0pt
* \topskip=11.0pt
* \footskip=30.0pt
* \marginparwidth=59.0pt
* \marginparsep=10.0pt
* \columnsep=10.0pt
* \skip\footins=10.0pt plus 4.0pt minus 2.0pt
* \hoffset=0.0pt
* \voffset=0.0pt
* \mag=1000
* \@twocolumnfalse
* \@twosidefalse
* \@mparswitchfalse
* \@reversemarginfalse
* (lin=72.27pt=25.4mm, 1cm=28.453pt)
```

```
(/usr/share/texlive/texmf-dist/tex/latex/ucs/ucsencs.def)
(/usr/share/texlive/texmf-dist/tex/latex/hyperref/nameref.sty
(/usr/share/texlive/texmf-dist/tex/generic/oberdiek/gettitlestring.sty))
```

Package hyperref Warning: Rerun to get /PageLabels entry.

```
(/usr/share/texlive/texmf-dist/tex/latex/psnfss/otlpppl.fd)
(/usr/share/texlive/texmf-dist/tex/latex/psnfss/omlzpplm.fd)
(/usr/share/texlive/texmf-dist/tex/latex/psnfss/omszpplm.fd)
(/usr/share/texlive/texmf-dist/tex/latex/psnfss/omxzpplm.fd)
(/usr/share/texlive/texmf-dist/tex/latex/psnfss/otlzpplm.fd)
(/usr/share/texlive/texmf-dist/tex/latex/jknapltx/ursfs.fd)
```

LaTeX Warning: No \author given.

(/usr/share/texlive/texmf-dist/tex/generic/oberdiek/se-ascii-print.def)

(/usr/share/texmf/tex/latex/lm/tl1lmtt.fd)

(/usr/share/texmf/tex/latex/lm/ts1lmtt.fd) [1] [2] [3]

Underfull \hbox (badness 10000) in paragraph at lines 550--551

[4]

Underfull \hbox (badness 10000) in paragraph at lines 577--578

[5]

! FancyVerb Error:

Empty verbatim environment

.
\FV@Error ...ncyVerb Error:^^J\space \space #1^^J}

1.649 ...60 [00:10<10:19, 10.50s/it]\end{Verbatim}

?

! Emergency stop.

\FV@Error ...ncyVerb Error:^^J\space \space #1^^J}

1.649 ...60 [00:10<10:19, 10.50s/it]\end{Verbatim}

Output written on notebook.pdf (5 pages).

Transcript written on notebook.log.

[NbConvertApp] PDF successfully created

[NbConvertApp] Writing 367982 bytes to Lab3Final.pdf

In []: